

# Sequence models

Attila Bagoly

1

# Applications

- Speech recognition
  - Input: sequence of pressure values
  - Output: sequence of words
- Music generation
  - Input: 0
  - Output sequence of notes
- Sentiment classification
  - Input: sequence of words
  - Output: rating (1-5)
- Machine translation
  - Input: sequence of words
  - Output: sequence of words
- Video activity recognition, summarization, etc.:
  - Input: sequence of pictures
  - Output: labels, sequence of words, etc.
- DNA sequence analysis:
  - Input: AGCCCT...
  - Output: eg. sequence of A,G,C,T which corresponds to a protein

# Representation

- **x**: Hello World

$$x^{<1>} \quad x^{<2>}$$

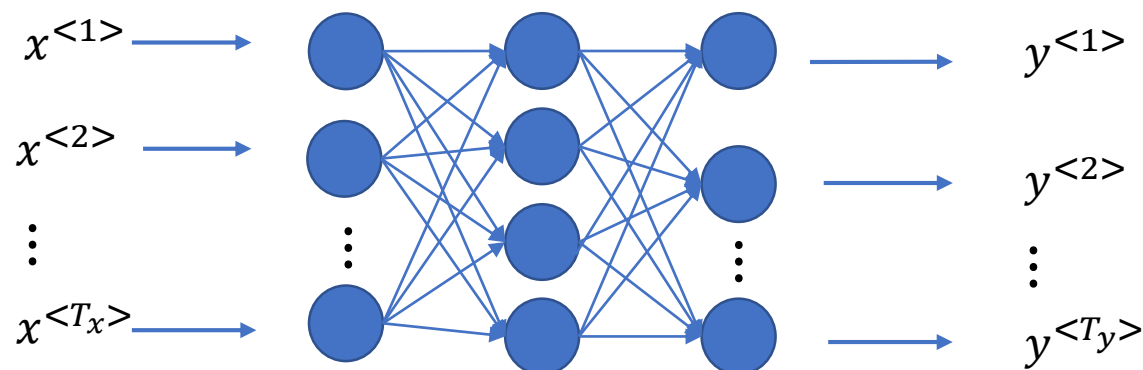
- Or character level:  $x^{<1>} = h$ ,  $x^{<2>} = e$ ,  $x^{<3>} = l \dots$

- Vocabulary: list of words or characters:

$$\begin{bmatrix} a \\ b \\ \vdots \end{bmatrix} \text{ or } \begin{bmatrix} a \\ and \\ \vdots \end{bmatrix}$$

- We encode the characters or words with the position index in the vocabulary (after pre-processing)
- Input: one-hot encoded indexes or it's lower dimensional embedding
- Song: vocabulary of notes

# Neural network

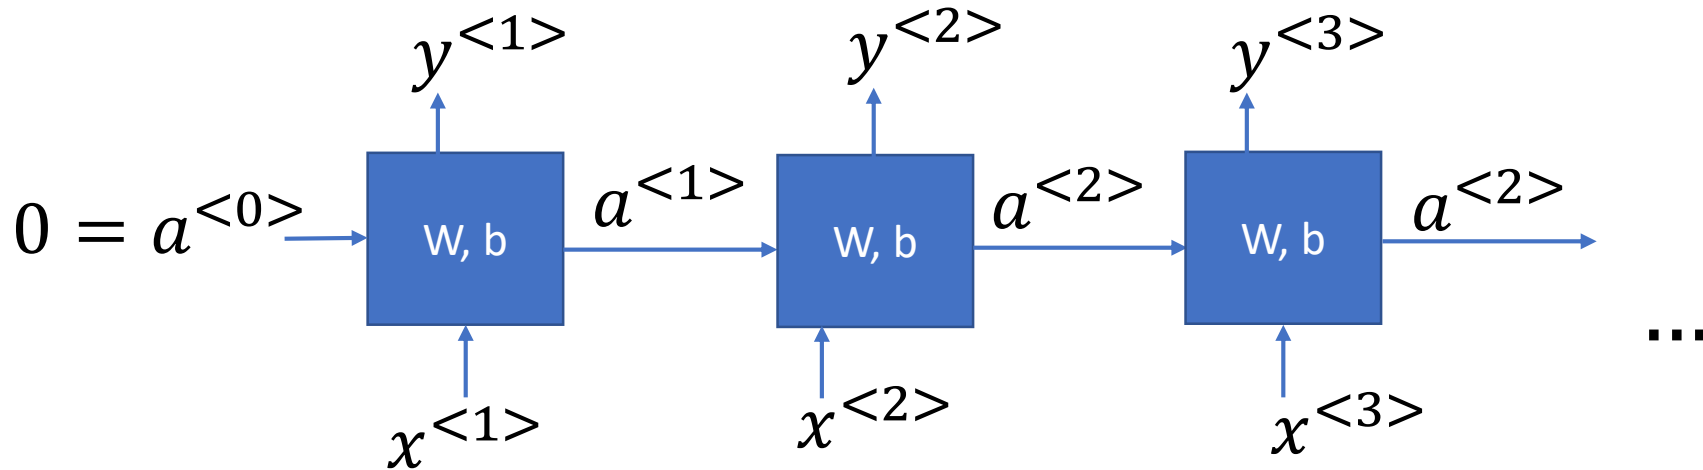


- Different input and output sizes:
  - Text sizes will vary example by example
  - Find the maximal text, pad with zeros the others (not too good)
- Bigger problem: doesn't share learned features across different positions of the text


Apple is good for your health. Apple juice is cheap.


You learned in position 1 that apple is a fruit (found some weights). In position 7 you can't use the weights from position 1, you would have to learn the same weights again.

# Recurrent neural networks (RNNs)



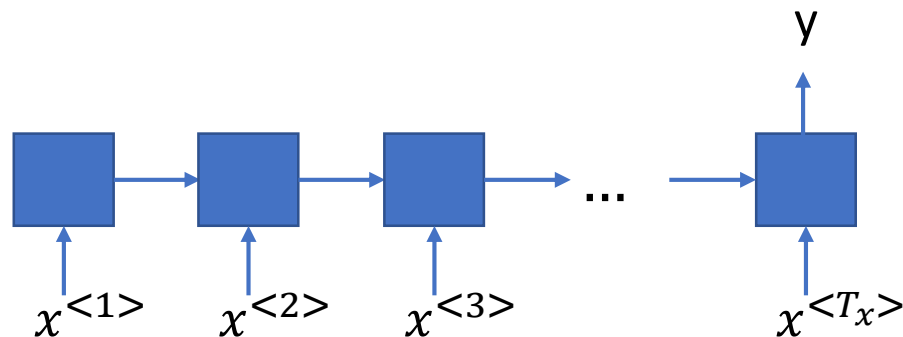
- $a^{<1>} = g(W_{aa}a^{<0>} + W_{ax}x^{<1>} + b_a)$
- $y^{<1>} = g(W_{ya}a^{<1>} + b_y)$
- $a^{<2>} = g(W_{aa}a^{<1>} + W_{ax}x^{<2>} + b_a)$
- $y^{<2>} = g(W_{ya}a^{<2>} + b_y)$
- ...



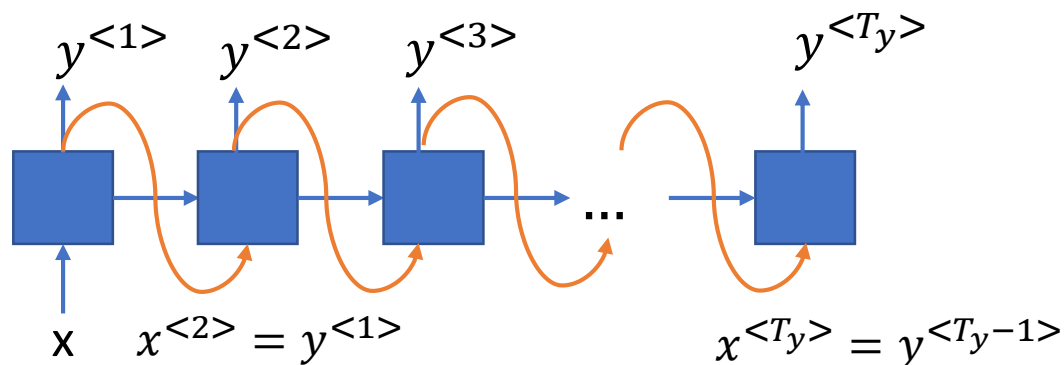


# RNN architectures

- Many-to-one (e.g. sentiment classification)



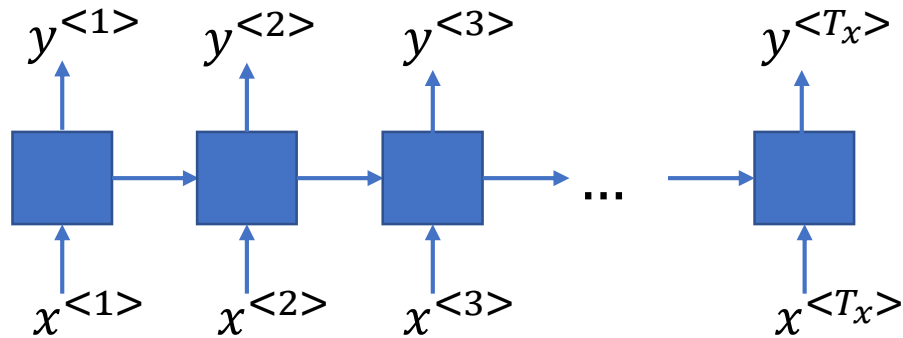


- One-to-many (e.g. music generation)



# RNN architectures

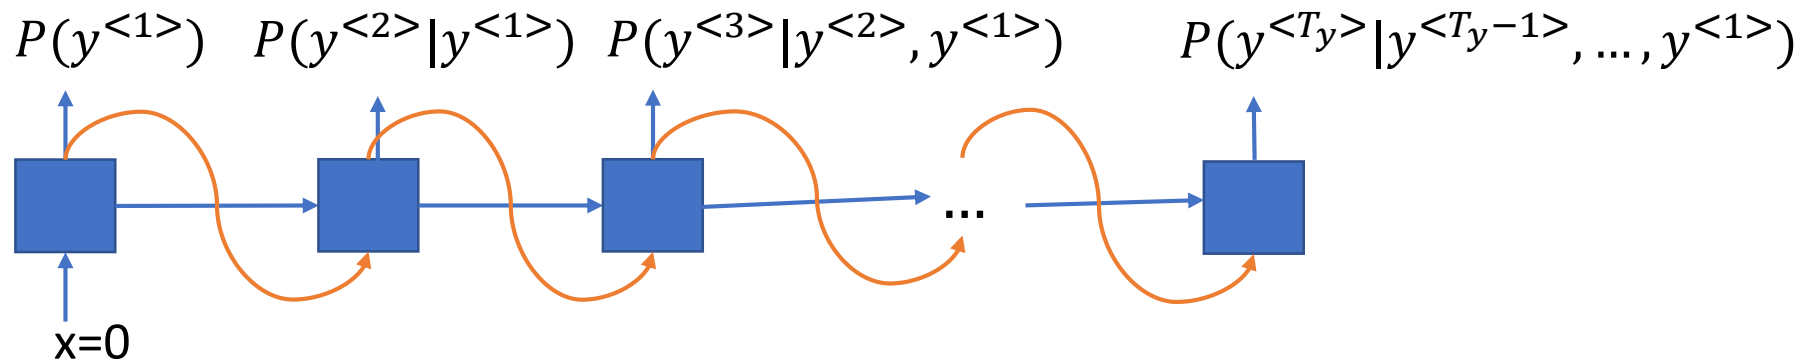
- Many-to-many: 2 case:  $T_x = T_y$  or  $T_x \neq T_y$





# Language modelling

- Language model:  $P(\text{sentences})$
- Language model with RNN:



$$P(y^{<1>}, y^{<2>}, \dots, y^{<T_y>}) =$$

$$P(y^{<1>}) P(y^{<2>} | y^{<1>}) P(y^{<3>} | y^{<2>} y^{<1>}) \dots P(y^{<T_y>} | y^{<T_y-1>}, \dots)$$

# Backpropagation


- Loss function:

$$L^{<t>}(\hat{y}^{<t>}, y^{<t>}) = - \sum_k y_k^{<t>} \log \hat{y}_k^{<t>}$$

$$L = \sum_{t=1}^{T_y} L^{<t>}(\hat{y}^{<t>}, y^{<t>})$$

- RNN:  $a^{<t>} = g(W_{aa}a^{<t-1>} + W_{ax}x^{<t-1>} + b_a)$

- Let's say we know:  $\frac{\partial L}{\partial a^{<t>}} = \frac{\partial L}{\partial y^{<t>}} \frac{\partial y^{<t>}}{\partial a^{<t>}}$

- We need:  $\frac{\partial L}{\partial W_{aa}}, \frac{\partial L}{\partial W_{ax}}, \frac{\partial L}{\partial b_a}$    $\frac{\partial L^{<t>}}{\partial W_{aa}} = \frac{\partial L^{<t>}}{\partial a^{<t>}} \frac{\partial a^{<t>}}{\partial W_{aa}}$

$$\frac{\partial a^{<t>}}{\partial W_{aa}} = g'(a^{<t>}) \cdot \left( a^{<t-1>} + W_{aa} \frac{\partial L}{\partial a^{<t-1>}} \right) \img alt="blue arrow pointing right" data-bbox="828 778 891 818"/> \text{Backprop through time}$$

- When we train we have to do backpropagation through time
- Long sentences: vanishing gradient problem
- Vanishing gradient:  $g'$  small,  $(g')^N \rightarrow 0$  (long sentence = big  $N$ )
- Similar to very deep networks
- ResNet: skip connection (linear information flow)

$$\begin{aligned} F(x) &\rightarrow F(x) + x \\ F'(x) \frac{\partial L}{\partial x} &\rightarrow F'(x) \frac{\partial L}{\partial x} + \frac{\partial L}{\partial x} \end{aligned}$$

- NEW architectures instead of RNN with linear information flow

# Gate recurrent unit (GRU)

- New hidden variable:  $c$  = memory cell
- Candidate value:  $\tilde{c} = g(W_{cc}c^{<t-1>} + W_{cx}x^{<t>} + b_c)$
- Update gate:  $\Gamma_u = \sigma(W_{uc}c^{<t-1>} + W_{ux}x^{<t>} + b_u)$
- Update rule:  $c^{<t>} = \Gamma_u \cdot \tilde{c} + (1 - \Gamma_u) \cdot c^{<t-1>}$
- Activation: simply  $a^{<t>} = c^{<t>}$
- Example:  $c_0$  = plural (0) / singular (1)


$$\begin{array}{ccc} \Gamma_u = 0 & \Gamma_u = 1 & \Gamma_u = 0 \\ c_0^{<1>} = 0 & c_0^{<2>} = 1 & c_0^{<3>} = 1 \end{array}$$

$$\begin{array}{cc} \Gamma_u = 0 & \Gamma_u = 0 \\ c_0^{<7>} = 1 & c_0^{<8>} = 1 \end{array}$$

$$\begin{array}{c} \Gamma_u = 1 \\ c_0^{<12>} = 0 \end{array}$$

The cat, who belongs to my friend, was red. The other cats ...

# Gate recurrent unit (GRU)




# Long short-term memory (LSTM)

- GRU is a simpler model which is able to learn long-term dependencies
- More powerful model is LSTM
- Solves the vanishing problem very similarly: linear information flow through time
- Candidate value:  $\tilde{c} = g(W_{cc}c^{<t-1>} + W_{cx}x^{<t>} + b_c)$
- It uses more gates:
  - Update gate:  $\Gamma_u = \sigma(W_{uc}c^{<t-1>} + W_{ux}x^{<t>} + b_u)$
  - Forget gate:  $\Gamma_f = \sigma(W_{fc}c^{<t-1>} + W_{fx}x^{<t>} + b_f)$
  - Output gate:  $\Gamma_o = \sigma(W_{oc}c^{<t-1>} + W_{ox}x^{<t>} + b_o)$
- Update rule:  $c^{<t>} = \Gamma_u \cdot \tilde{c} + \Gamma_f \cdot c^{<t-1>}$
- Activation:  $a^{<t>} = \Gamma_o \cdot g(c^{<t>})$

# Long short-term memory (LSTM)

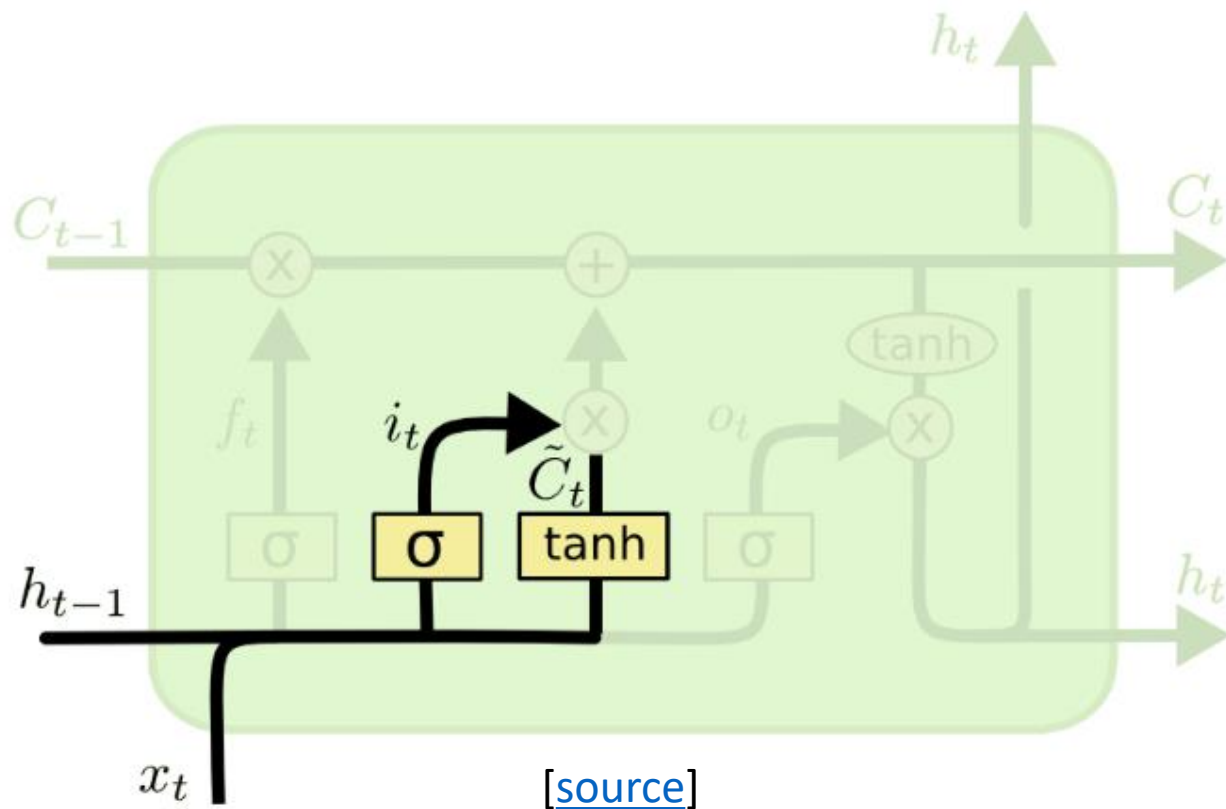
- Forget gate:
  - New parameters which control what information we forget



[[source](#)]

# Long short-term memory (LSTM)

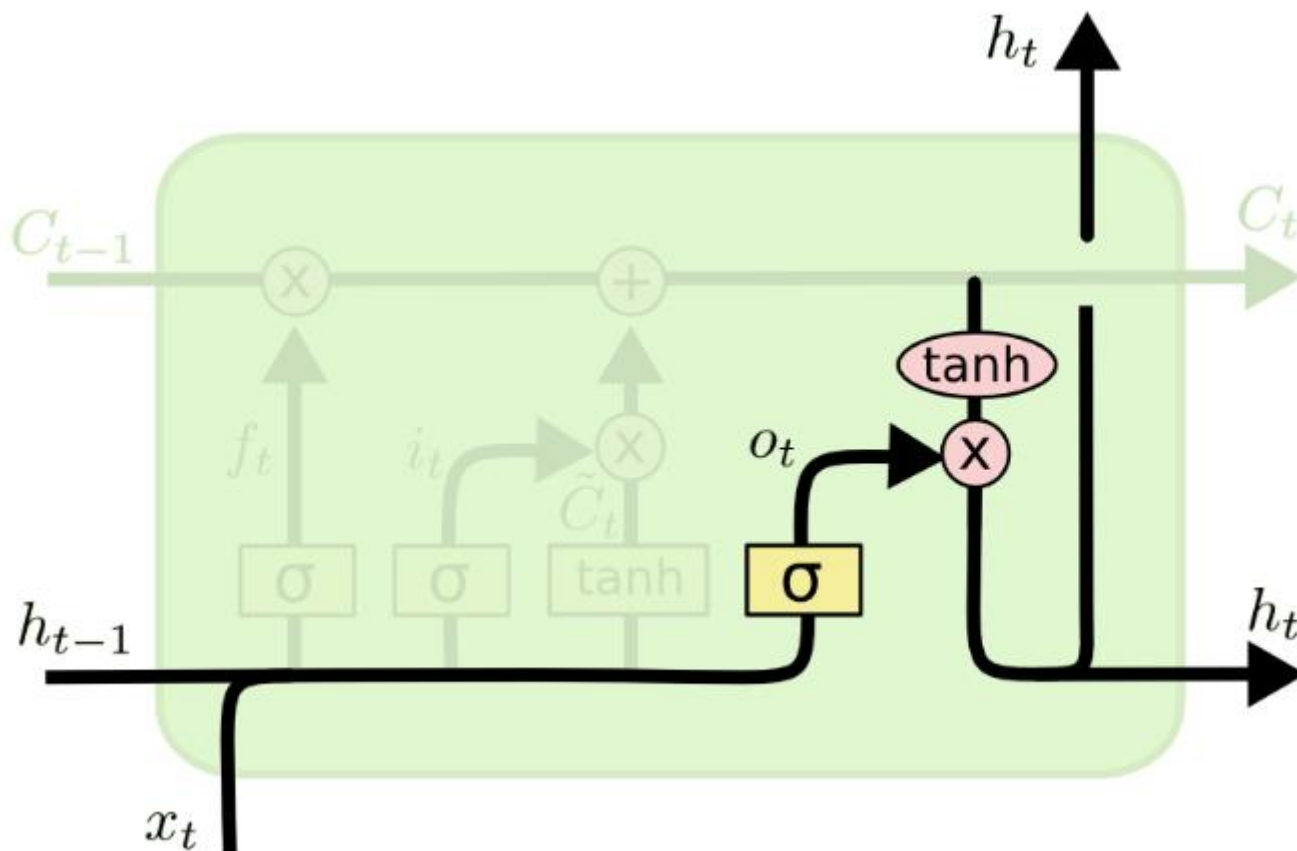
- We learn parameters for a candidate value: new  $c$  vector
- We also learn parameters for an update gate: how much information we want to store





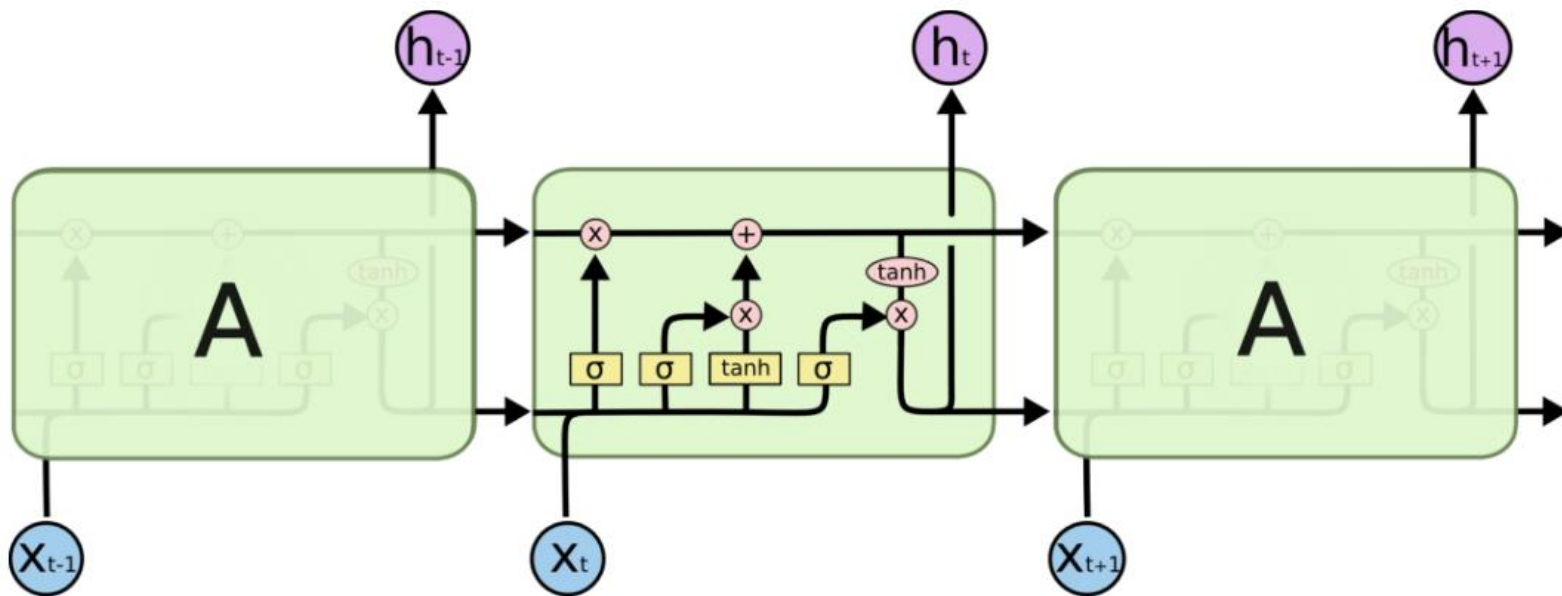
# Long short-term memory (LSTM)

- Output gate: what information we want to the cell to output
- This gate also has learned parameters




[[source](#)]

- The full LSTM cell



[[source](#)]

- We can stack RNN, GRU or LSTM cells in top of each other
- Typically a few cell (e.g. LSTM contains 4 “layers” => few layer is a deep network)



# Demo notebooks

- Music generation:

[https://github.com/qati/DeepLearningCourse/blob/master/demo\\_notebooks/lecture\\_09/piano.ipynb](https://github.com/qati/DeepLearningCourse/blob/master/demo_notebooks/lecture_09/piano.ipynb)

- Tweet generation:

[https://github.com/qati/DeepLearningCourse/blob/master/demo\\_notebooks/lecture\\_09/tweets.ipynb](https://github.com/qati/DeepLearningCourse/blob/master/demo_notebooks/lecture_09/tweets.ipynb)