


Autoencoders, VAE, GAN

Attila Bagoly

1


Autoencoders

- Unsupervised learning: no-labels (just X)
- Autoencoder: $\mathcal{M}: X \rightarrow X$
- Tries to learn “identity” function, but we add **constraints**
- **Constraints:** e.g. less neurons in the middle (lower dim. repr.)
- Design:
 - Encoder: encodes the input into a lower dimensional space
 - Decoder: decodes the input from the lower dimensional space to the original space
- PCA: linear projection; Autoencoder: non-linear projection




Autoencoders

- Loss: e.g. MSE, binary-crossentropy (minimizing the difference between the input and output)
- Non-triviality: lower dimensional representation



Convolutional autoencoder

- The encoder is a classical convolutional network
- Decrease the spatial dimension:
 - MaxPooling, AvgPooling: not learnable
 - Convolution with stride >1: learnable
(we learn how to do the downsampling, tiny details, preserves spatial information)
 - Size calculation: $d' = \frac{d-f+2p}{s} + 1$
 - Padding: 'valid': $p=0$; 'same': $d = d' \rightarrow p$
- What about in the decoder?
- We need to reverse Pooling and Conv
- Encoder: downsampling
- Decoder: upsampling



Upsampling: “max unpooling”

- Pooling: remember the max positions
- Unpooling: lot of zeros, except in the max positions
- Fixed layer: doesn't learn

| | | | |
|------------|------------|------------|------|
| 0.1 | 0.5 | 1.2 | -0.7 |
| 0.8 | -0.2 | -0.5 | 0.3 |
| 0.4 | 0.9 | -0.1 | -0.2 |
| -0.6 | 0.1 | 0.5 | 0.3 |

max-pooling

| | |
|-----|-----|
| 0.8 | 1.2 |
| 0.9 | 0.5 |

| | | | |
|---|---|---|--|
| | | x | |
| x | | | |
| | x | | |
| | | x | |

max locations


| | | | |
|------------|------------|------------|---|
| 0 | 0 | 0.5 | 0 |
| 1.3 | 0 | 0 | 0 |
| 0 | 0.4 | 0 | 0 |
| 0 | 0 | 0.1 | 0 |

unpooling

| | |
|-----|-----|
| 1.3 | 0.5 |
| 0.4 | 0.1 |


Upsampling: Transposed convolution

- Other name: Fractionally strided convolution
- Sometimes called: deconvolution
- But deconvolution (inverse convolution) exists and it is mathematically very different: but both results the same dimension. Real deconv not used in deep learning!
- Transposed convolution: learnable layer (learns how to do the upsampling)





Upsampling: Transposed convolution

- Transposed convolution is also a convolution
- But with some fancy padding




Convolutional autoencoder

- Decoder: upsample with 'max unpooling' or with transposed convolution



Denoising autoencoder




Original
Image

Noisy
Input


Code

Output


Original Images



Noisy Input



Autoencoder Output




Sparse autoencoder

- Average activation in 2. layer:

$$\hat{\rho}_j = \frac{1}{m} \sum_{i=1}^m \left[a_j^{(2)}(x^{(i)}) \right]$$



- Sparsity parameter: ρ
- Constraint: $\hat{\rho}_j = \rho$
- We want: average activation of the hidden unit to be close to: ρ
- Constraint: loss penalty term



$$\sum_{j=1}^{s_2} \rho \log \frac{\rho}{\hat{\rho}_j} + (1 - \rho) \log \frac{1 - \rho}{1 - \hat{\rho}_j} = \sum_{j=1}^{s_2} \text{KL}(\rho || \hat{\rho}_j)$$

Autoencoder applications


- Pretraining
- Data compression: hashing
- Image search
- Information retrieval
- Denoising, reconstruction
- Image colorization
- Generating higher resolution images



Autoencoder demo notebooks

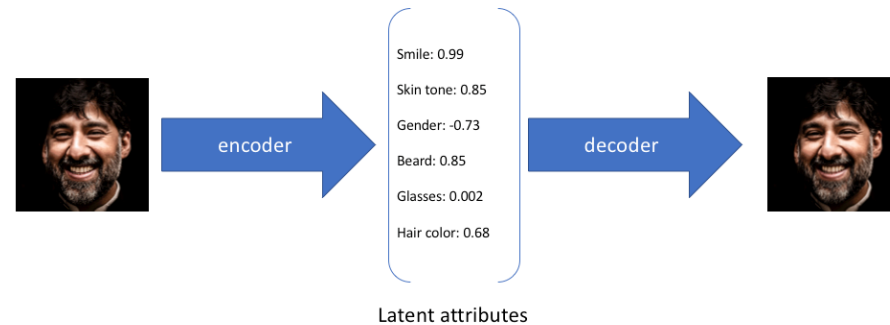
Problems with autoencoder

- Latent space not continuous
- Not well structured: hard to interpret it
- Clusters: easier to decode
- Gaps: if we change the latent code a little bit, we can get totally unrealistic images
- Happy, not happy encoded -> what if we move in “happy direction”?
- We can fall into gaps
- Can't generate new images





Variational autencoder




- Ideal autencoder:




- **VAE**: continuous, structured latent space by design
- Each latent variable range of possible values
- Each latent var for given input: prob. dist



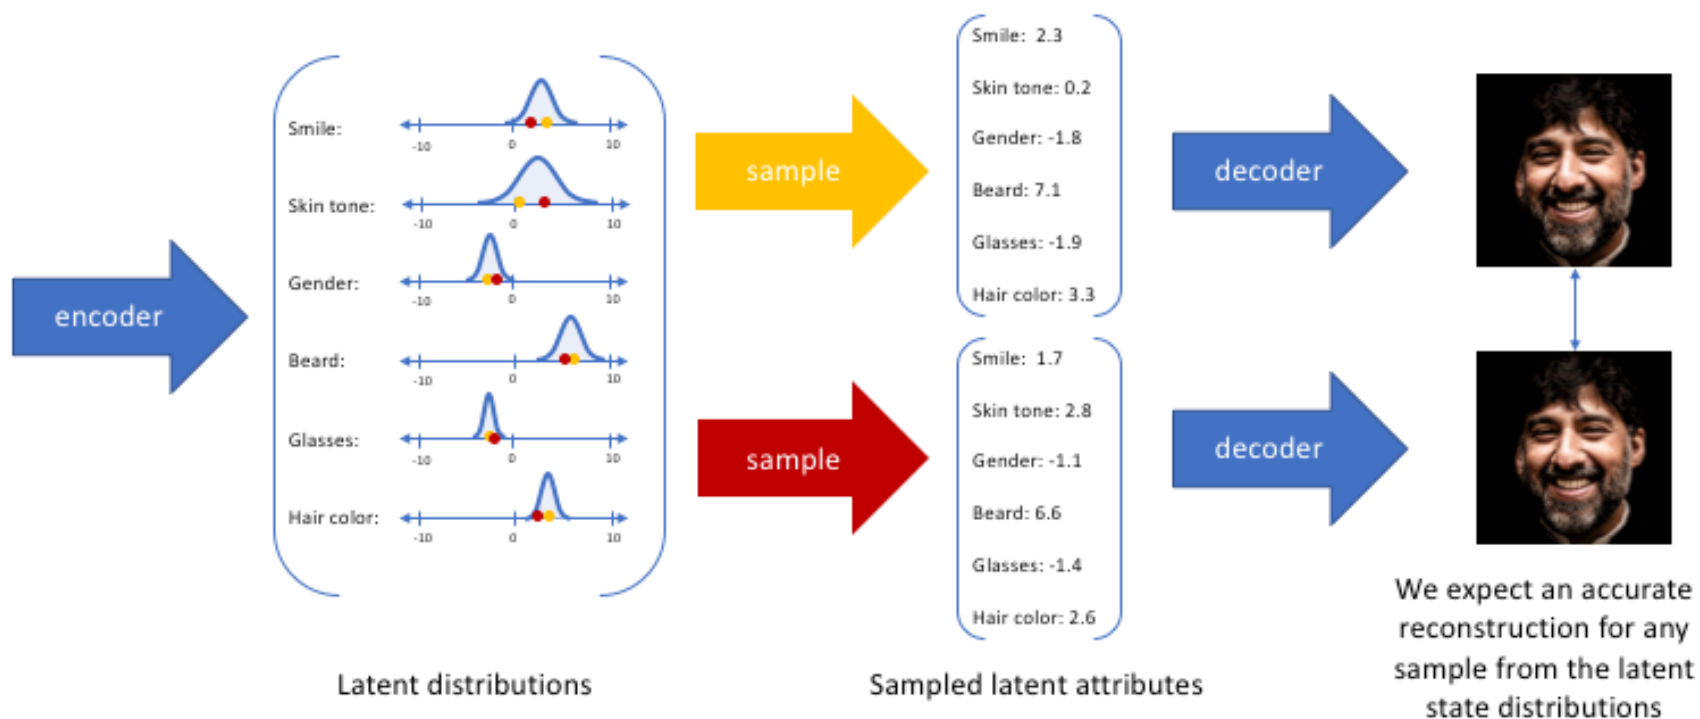
vs.



- Given input: each latent variable is a probability distribution
- Encoder: encodes the input to prob. distributions
- Decoder: decodes randomly sampled z from these distrs.




- Encoder: range of values; Decoder: random sample
➡ smooth latent representation
- Nearby in latent space: very similar reconstruction



- We have latent variable z , and observation x
- Latent variable: $z \sim p(z)$ (prior)
- Draw datapoint: $x \sim p(x|z)$ (likelihood)
- We want to learn what is z given x
- In other words: $p(z|x)$:

$$p(z|x) = \frac{p(x|z)p(z)}{p(x)} = \frac{p(x|z)p(z)}{\int p(x|z)p(z)dz}$$

- Problem: integral intractable
- Solution: variational inference: approximate posterior with $q(z|x)$, which has tractable distribution
- We choose the parameters of q to be a good approximate




- We need $p(z|x)$, but we can't calculate it
- We approximate it with $q(z|x)$
- We want to $q(z|x)$ to be close to $p(z|x)$:

$$\min \text{KL}(q(z|x)|p(z|x))$$

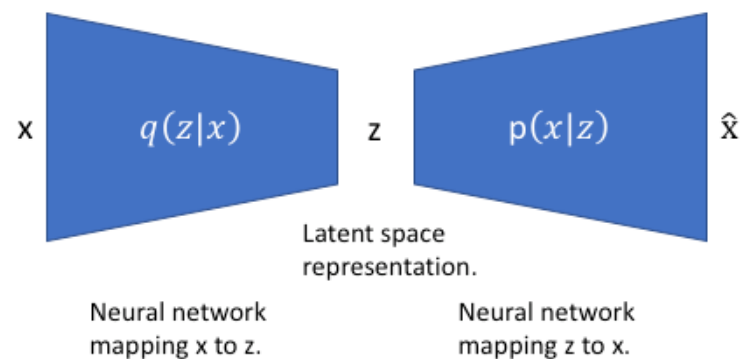
- After some calculation we will get:

$$\mathbb{E}_{q(z|x)} \log p(x|z) - \text{KL}(q(z|x)|p(z))$$

- First term:
reconstruction
error
- Second:
we want $q(z|x)$
to be close to
the prior



We'd like to use our observations to understand the hidden variable.



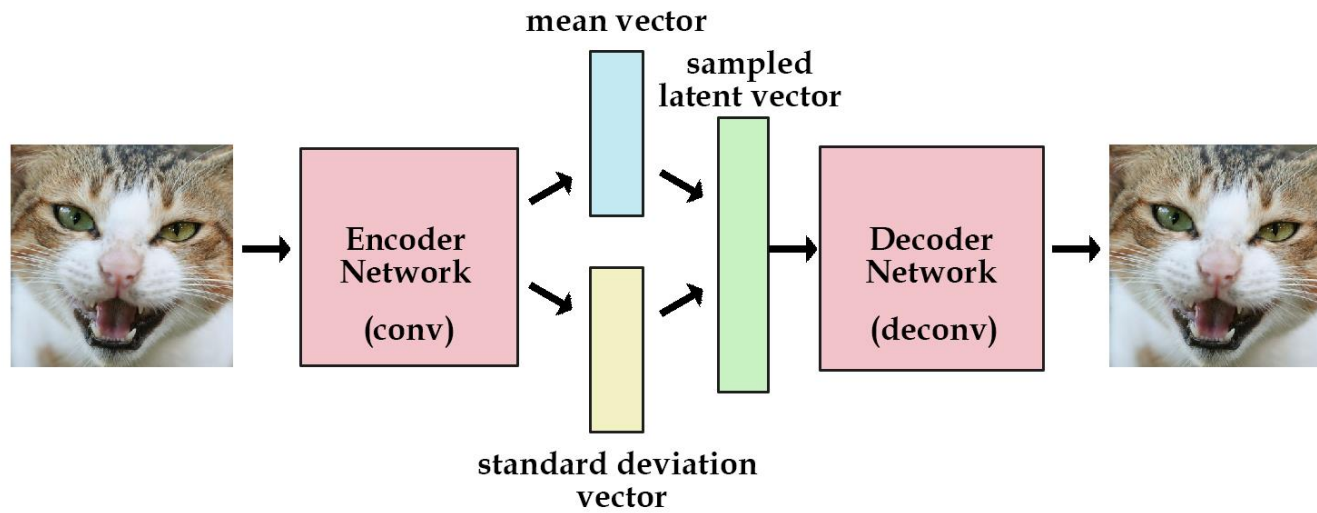
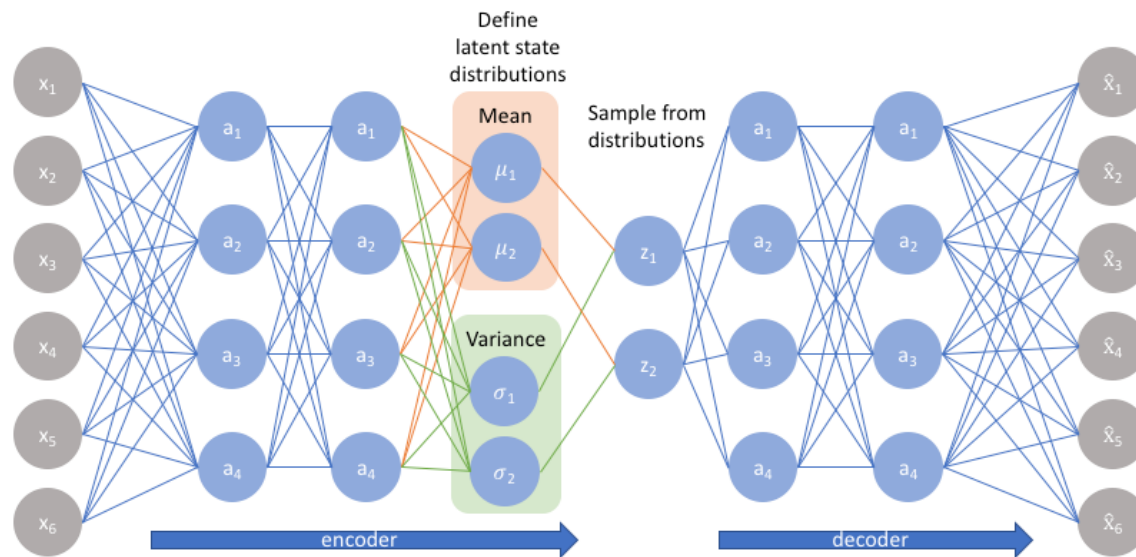
- Autoencoder: $\mathcal{L}(\hat{x}, x)$ (gaps in latent code z)
- Variational autoencoder (\sum each dim. in latent space):

$$\mathcal{L}(\hat{x}, x) + \sum_j \text{KL} \left(q_j(z|x) | p(z) \right)$$

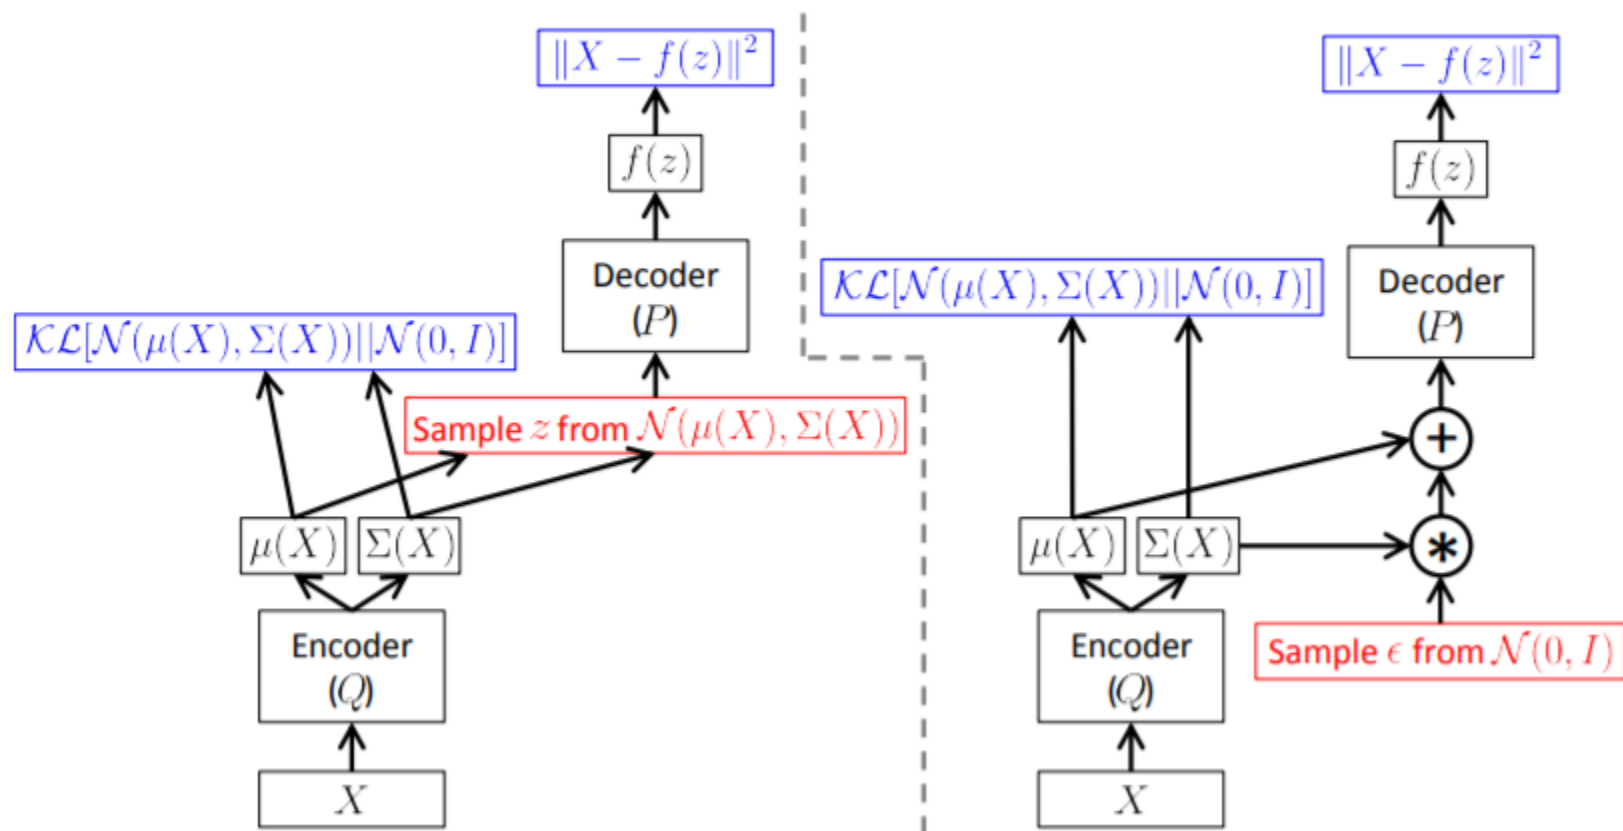
- Prior: $p(z) \sim \mathcal{N}(0, I)$
- q : $q(z|x) \sim \mathcal{N}(\mu(x), \Sigma(x))$
- Kullback-Leibler divergence (k dimension):

$$\begin{aligned} & \text{KL} \left(\mathcal{N}(\mu(x), \Sigma(x)) | \mathcal{N}(0, I) \right) \\ &= \frac{1}{2} \left(\text{tr} \Sigma(x) + (\mu(x))^T \mu(x) - k - \log \det \Sigma(x) \right) \end{aligned}$$

VAE architecture




VAE architecture




VAE latent space

Penalizing reconstruction loss encourages the distribution to describe the input



Our distribution deviates from the prior to describe some characteristic of the data


Without regularization, our network can “cheat” by learning narrow distributions



With a small enough variance, this distribution is effectively only representing a single value


Penalizing KL divergence acts as a regularizing force

Attract distribution to have zero mean




Ensure sufficient variance to yield a smooth latent space


Only reconstruction loss




Only KL divergence



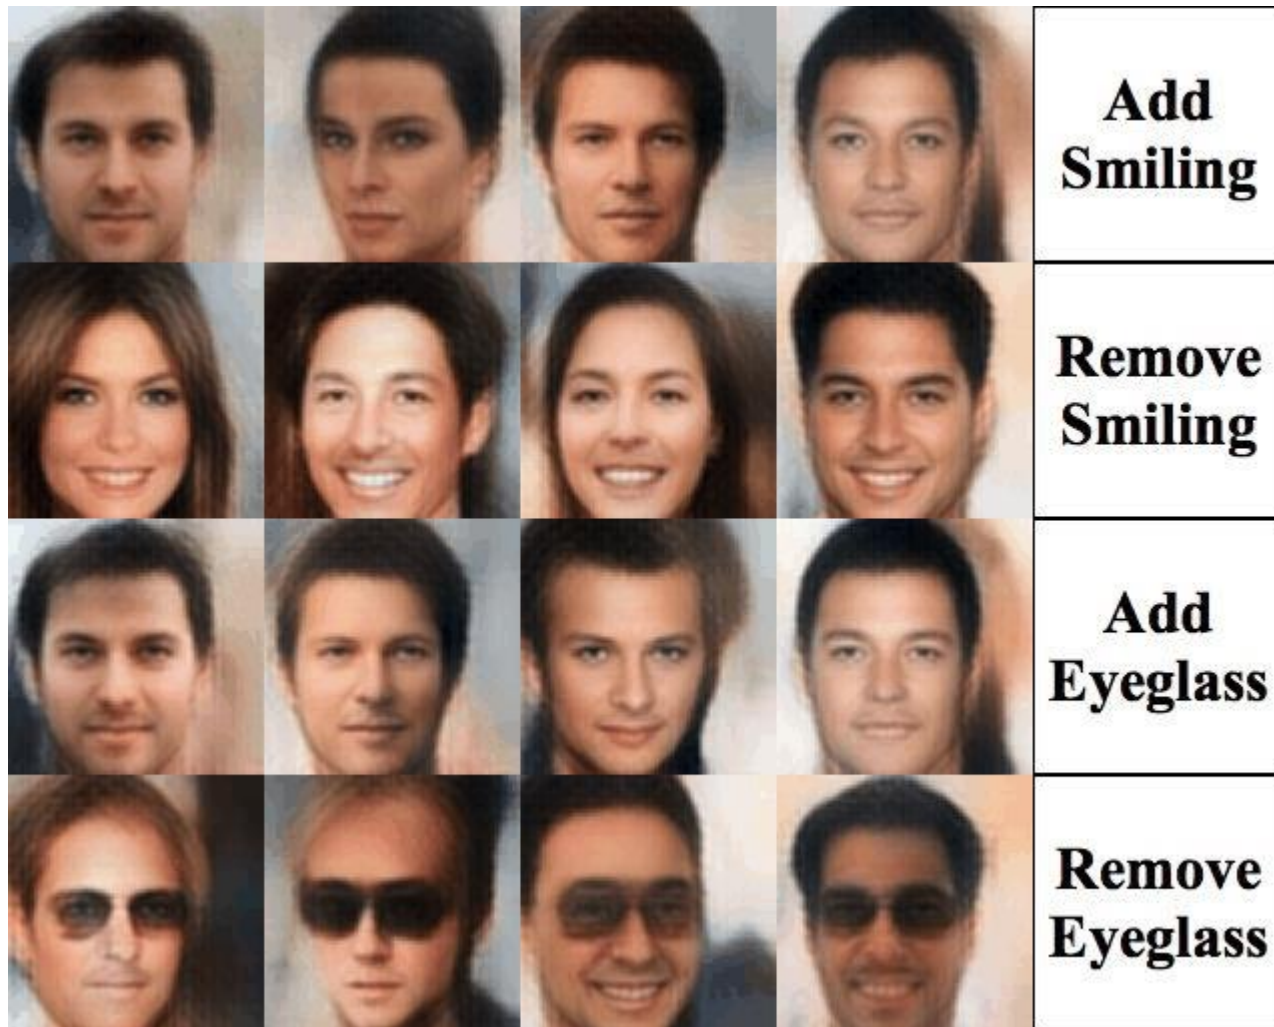
Combination





- **DEMO notebooks**

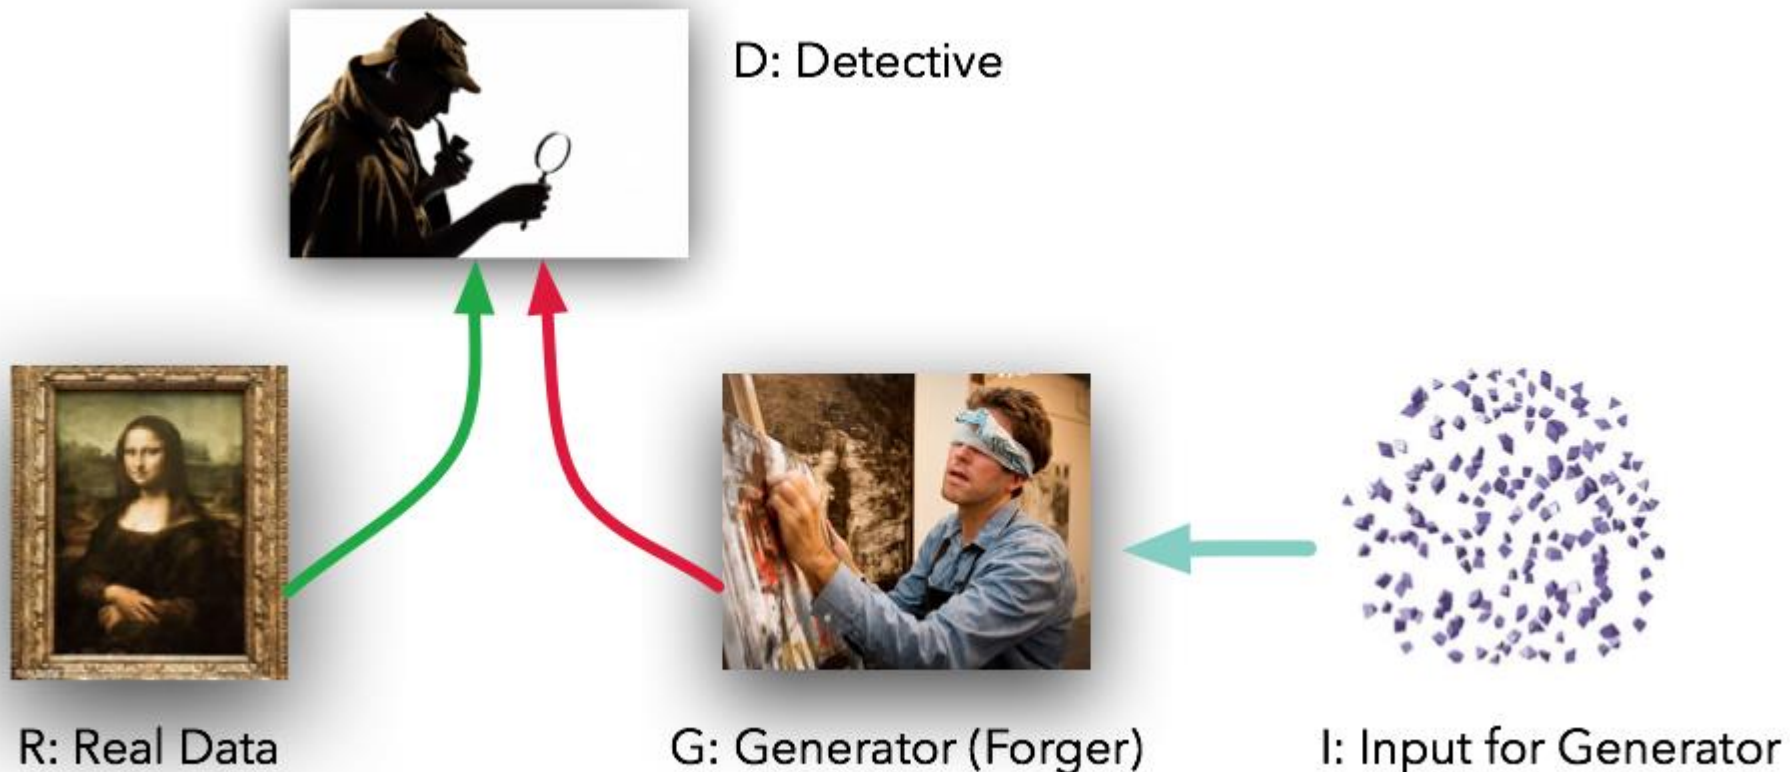
VAE example



- <https://www.youtube.com/watch?v=Q1XuXwPVFko>
- <https://magenta.tensorflow.org/music-vae>

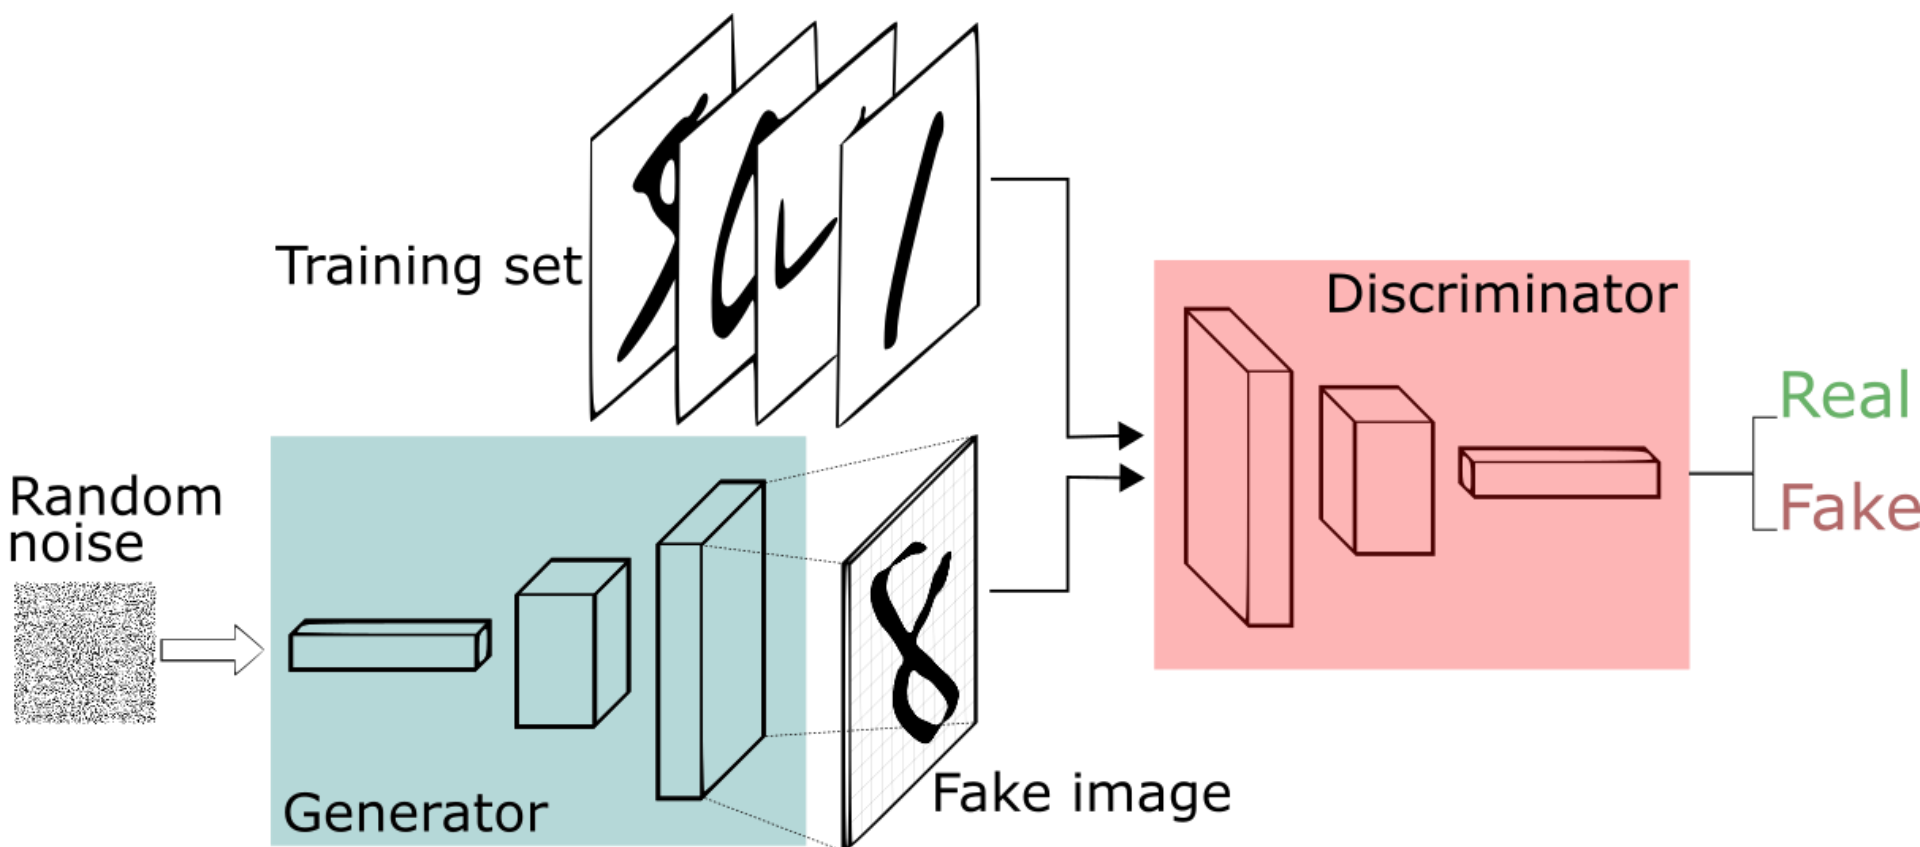
Generative adversarial networks (GAN)

- Imagine an art forger who wants to make Mona Lisa to sell it
- But the gallery has an art “detective”: game



Generative adversarial networks (GAN)

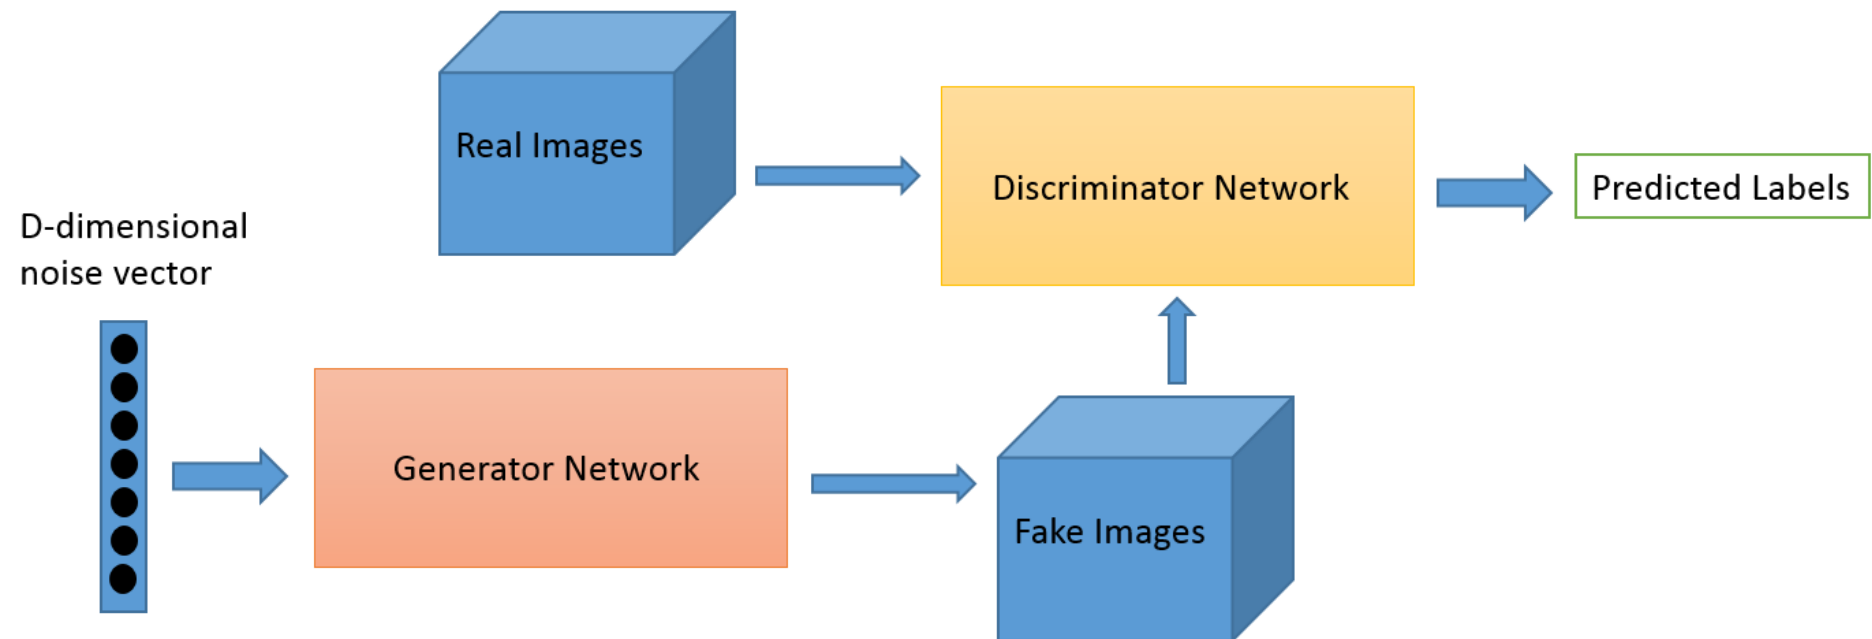
- GAN: two competing network
- Generator: tries to create real-like pictures
- Discriminator: wants to detect the forgery



- The cost used for the discriminator:

$$J^{(D)}(\Theta^{(D)}, \Theta^{(G)}) = -\frac{1}{2} \mathbb{E}_{x \sim p(x)} \log D(x) - \frac{1}{2} \mathbb{E}_{x \sim p(x)} \log (1 - D(G(x)))$$

- This is the binary-crossentropy (real=1, fake=0)



- Discriminator tries to distinguish between real and fake data
- Generator tries to fool the discriminator
- What loss can be used for this?

- Simplest case:

$$J^{(G)}(\Theta^{(D)}, \Theta^{(G)}) = -J^{(D)}(\Theta^{(D)}, \Theta^{(G)}) = \frac{1}{2} \mathbb{E}_{x \sim p(x)} \log D(x) + \frac{1}{2} \mathbb{E}_{x \sim p(x)} \log (1 - D(G(x)))$$

- So the discriminator tries to make $D(G(x))$ close to 0
- The generator tries to make $D(G(x))$ close to 1 (log big negative number)

- Minimax game:

$$\min_G \max_D V(D, G)$$

$$V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))]$$

Problems with Counting




(Goodfellow 2016)

Problems with Perspective



(Goodfellow 2016)

Problems with Global Structure



(Goodfellow 2016)

Examples

- Text to image: <https://arxiv.org/pdf/1703.06412.pdf>

This flower has yellow petals along with green and yellow stamen



This flower is red and yellow in color, with petals that are ruffled and curled



This flower has petals that are yellow with red lines



This flower is white and pink in color, with petals that are oval shaped



A yellow flower with large petal with a large long pollen tubes



The petals on this flower are white with yellow stamen



Examples

- Houses:

<https://www.youtube.com/watch?v=JCEuwO5BPnk&t=97s>

- Zebras:

<https://www.youtube.com/watch?v=9reHvktowLY>

- Faces:

<https://www.youtube.com/watch?v=G06dEcZ-QTg>

- And a lot of other amazing examples

- Hacks: <https://github.com/soumith/ganhacks>

- **DEMO notebooks**

Értékelés:

Házik:

- hw01 2 pont
- hw02 5.5 pont
- hw03 8 pont
- hw09 3 pont

Photoz kaggle:






- 1-10: 7 pont
- 11-20: 5 pont
- 21-baseline: 3 pont
- baseline alatt: 1 pont


Photoz kaggle:

- 1-10: 11 pont
- 11-20: 9 pont
- 21-baseline: 7 pont
- baseline alatt: 2 pont

Max pont 36.5


Max házi + (kaggle baseline + ϵ) = 28.5





Counter({1: 3, 2: 3, 3: 5, 4: 5, 5: 26})

Grades



| | githubName | hw01 | hw02 | hw03 | hw09 | happy_points | photoz_points | points_sum | grades |
|----|---------------|-------|------|------|------|--------------|---------------|------------|--------|
| 16 | danielgrajzel | 1.835 | 0.0 | 0.0 | 0.0 | 0.0 | 3.0 | 4.835 | 1 |
| 19 | gichy | 2.000 | 3.0 | 0.0 | 0.0 | 0.0 | 0.0 | 5.000 | 1 |
| 45 | Zongi93 | 0.000 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000 | 1 |

| | githubName | hw01 | hw02 | hw03 | hw09 | happy_points | photoz_points | points_sum | grades |
|----|------------|-------|------|-------|------|--------------|---------------|------------|--------|
| 10 | balint225 | 1.835 | 5.5 | 3.000 | 0.0 | 0.0 | 0.0 | 10.335 | 2 |
| 17 | e-velin | 2.000 | 0.0 | 0.000 | 0.0 | 9.0 | 0.0 | 11.000 | 2 |
| 18 | ggalgoczi | 2.000 | 2.0 | 2.165 | 0.0 | 0.0 | 7.0 | 13.165 | 2 |

| | githubName | hw01 | hw02 | hw03 | hw09 | happy_points | photoz_points | points_sum | grades |
|----|------------|-------|------|-------|------|--------------|---------------|------------|--------|
| 1 | CliffyH | 2.000 | 5.5 | 6.000 | 3.0 | 0.0 | 3.0 | 19.500 | 3 |
| 24 | ilxstatus | 2.000 | 5.5 | 4.000 | 2.5 | 2.0 | 3.0 | 19.000 | 3 |
| 26 | kazozoka | 2.000 | 5.0 | 6.000 | 0.0 | 0.0 | 5.0 | 18.000 | 3 |
| 27 | kissmate6 | 1.835 | 5.0 | 3.000 | 2.5 | 0.0 | 5.0 | 17.335 | 3 |
| 37 | oresme | 2.000 | 3.5 | 3.665 | 2.5 | 7.0 | 0.0 | 18.665 | 3 |

| | githubName | hw01 | hw02 | hw03 | hw09 | happy_points | photoz_points | points_sum | grades |
|----|-------------|-------|------|------|------|--------------|---------------|------------|--------|
| 4 | PentadD | 2.000 | 5.5 | 3.0 | 2.5 | 7.0 | 3.0 | 23.000 | 4 |
| 6 | Turcsi | 1.835 | 5.0 | 6.0 | 0.0 | 7.0 | 3.0 | 22.835 | 4 |
| 28 | kommancs96 | 1.835 | 5.0 | 8.0 | 3.0 | 0.0 | 5.0 | 22.835 | 4 |
| 33 | masterdesky | 1.500 | 5.0 | 6.0 | 0.0 | 7.0 | 1.0 | 20.500 | 4 |
| 41 | zentaijanos | 2.000 | 4.0 | 1.5 | 3.0 | 9.0 | 3.0 | 22.500 | 4 |

Homework for extra points

- <https://github.com/qati/DeepLearningCourse/tree/master/assignments/emoji>

References

- <https://towardsdatascience.com/applied-deep-learning-part-3-autoencoders-1c083af4d798>
- <https://towardsdatascience.com/types-of-convolutions-in-deep-learning-717013397f4d>
- David, Eli & Netanyahu, Nathan. (2016). DeepPainter: Painter Classification Using Deep Convolutional Autoencoders. 20-28. 10.1007/978-3-319-44781-0_3.
- Stanford CS231n, Lecture 11: <https://www.youtube.com/watch?v=nDPWywWRIRo>
- <https://towardsdatascience.com/autoencoders-introduction-and-implementation-3f40483b0a85>
- <https://blog.manash.me/implementing-pca-feedforward-and-convolutional-autoencoders-and-using-it-for-image-reconstruction-8ee44198ea55>
- <http://www.ericklwilkinson.com/blog/2014/11/19/deep-learning-sparse-autoencoders>
- <http://ufldl.stanford.edu/tutorial/unsupervised/Autoencoders/>
- <https://towardsdatascience.com/autoencoders-are-essential-in-deep-neural-nets-f0365b2d1d7c>
- <https://towardsdatascience.com/intuitively-understanding-variational-autoencoders-1bfe67eb5daf>
- <https://www.jeremyjordan.me/variational-autoencoders/>
- <https://jaan.io/what-is-variational-autoencoder-vae-tutorial/>
- <http://kvfrans.com/variational-autoencoders-explained/>
- <https://github.com/houxianxu/DFC-VAE>
- <https://deeplearning4j.org/generative-adversarial-network>
- <https://www.analyticsvidhya.com/blog/2017/06/introductory-generative-adversarial-networks-gans/>
- <https://medium.com/@devnag/generative-adversarial-networks-gans-in-50-lines-of-code-pytorch-e81b79659e3f>
- <https://arxiv.org/pdf/1701.00160.pdf>
- <https://github.com/hjweide/adversarial-autoencoder>
- Deep Learning with Python, by FRANÇOIS CHOLLET